

Databases

Storing Data

Format	Where
Fixed field record	Used in punch cards, each record can store 80 characters, some of the 80 characters are allocated to different fields.
Comma separated value record	More flexible than fixed field records, as there is no character limit to the value of a field.

A **simple associative store** will have the **API formal specification**.

```
store: string * string -> unit
retrieve: string -> string option (* either the value stored, or nothing at all *)
```

Definitions

Keyword	Meaning
Value	What is being stored.
Field	A place to hold a value.
Schema	A strongly typed specification of the database.
Key	The field(s) used to locate a record.
Index	A derived structure used to quickly find relevant record.
Query	A lookup function.
Update	A modification of data.
Transaction	An atomic change of a set of field.

An update is often implemented as a transaction.

Abstraction and Interfaces

We hope to have a fairly narrow interface that can map onto many operations. In an ideal world:

- All operations are done through the interface.
- The interface never change.

But in real world, large database projects are often a mess, and fixing the mess sometimes requires changing the interface. Changing interfaces break applications!

Logical Arrangement of a Database

The **DBMS** uses the **disc drive**, which is an abstraction over the **secondary storage**. This gives a consistent, nonchanging API which applications can use.

This API is often in form of SQL queries, SQL injection is an example of insecurity in the interface.

Operating System View of a Database

The **DBMS** is a process running in userspace.

- An **application** communicates with the **DBMS** through **kernel space communication**.
- Then the **DBMS** use **kernel space drivers** to access the **DBMS**.

An alternative implementation have the **DBMS** access a non-OS partition directly.

Data can be stored in primary store, secondary store, or distributed.

Definition

Big data is data too big to fit in primary store.

DBMS Operations

Operation	Description
Create	Insert new data
Read	Query the database
Update	Modify objects in the database
Delete	Remove data

Management Operations

- Create/change schema
- Create views
- Indexing/stats generation
- Reorganise data layout and backups

Amount of Writing

Database implementation choice depends on the amount of writing.

- A database can have a lot of lookups but rarely changes (e.g. library catalogue).
- **Transaction optimised**: concurrent queries and updates, they will affect the consistency of reads.

Definition

Atomicity is where changes are apparent to all users at once. An overhead is needed to achieve this.

- **Append only journal**: inverse of a transaction optimised DB, data is never updated.

Consistency Checks

Consistency rules includes value range check, foreign key referential integrity, value atomicity (all values are atomic), entity integrity (no missing fields).

Types of Databases**Relational Database**

Consists of 2D tables.

- One row per record (a.k.a. a tuple), each with a number of fields.
- A table can have a schema.
- The ordering of the fields are unimportant.

Distributed Databases

Database can be spread over multiple machines.

Benefit	Description
Scalability	Dataset may be too large for a single machine.
Fault tolerance	Service can survive the failure of some machines.
Lower latency	Data can be located closer to the users.

Downsides: overhead after an update to provide a consistent view.

Keyword	Meaning
Consistency	All reads return data that is up to date.
Accessibility	All clients can find some replica of the data.
Partition tolerance	The system continues to operate despite arbitrary message loss or failure of parts of the system.

It is impossible to achieve all 3 in a distributed database. (*why?*)

System that offers **eventual consistency** will eventually reach a consistent state once activity ceases.